

Primary configuration settings

#####

This configuration file is used to manage the behavior of the Salt Master.

Values that are commented out but have an empty line after the comment are

defaults that do not need to be set in the config. If there is no blank line

after the comment then the value is presented as an example and is not the

default.

Per default, the master will automatically include all config files

from master.d/*.conf (master.d is a directory in the same directory

as the main master config file).

#default_include: master.d/*.conf

The address of the interface to bind to:

#interface: 0.0.0.0

Whether the master should listen for IPv6 connections. If this is set to True,

the interface option must be adjusted, too. (For example: "interface: '::")

#ipv6: False

The tcp port used by the publisher:

#publish_port: 4505

The user under which the salt master will run. Salt will update all

permissions to allow the specified user to run the master. The exception is

the job cache, which must be deleted if this user is changed. If the

modified files cause conflicts, set verify_env to False.

#user: root

Tell the master to also use salt-ssh when running commands against minions.

#enable_ssh_minions: False

The port used by the communication interface. The ret (return) port is the

interface used for the file server, authentication, job returns, etc.

#ret_port: 4506

Specify the location of the daemon process ID file:

#pidfile: /var/run/salt-master.pid

The root directory prepended to these options: pki_dir, cachedir,

sock_dir, log_file, autosign_file, autoreject_file, extension_modules,

key_logfile, pidfile, autosign_grains_dir:

#root_dir: /

The path to the master's configuration file.

#conf_file: /etc/salt/master

Directory used to store public key data:

#pki_dir: /etc/salt/pki/master

Key cache. Increases master speed for large numbers of accepted

keys. Available options: 'sched'. (Updates on a fixed schedule.)

Note that enabling this feature means that minions will not be

available to target for up to the length of the maintenance loop

which by default is 60s.

#key_cache: "

Directory to store job and cache data:

This directory may contain sensitive data and should be protected accordingly.

#

#cachedir: /var/cache/salt/master

Directory for custom modules. This directory can contain subdirectories for

each of Salt's module types such as "runners", "output", "wheel", "modules",

"states", "returners", "engines", "utils", etc.

#extension_modules: /var/cache/salt/master/extmods

Directory for custom modules. This directory can contain subdirectories for

each of Salt's module types such as "runners", "output", "wheel", "modules",

"states", "returners", "engines", "utils", etc.

Like 'extension_modules' but can take an array of paths

#module_dirs: []

Verify and set permissions on configuration directories at startup:

#verify_env: True

Set the number of hours to keep old job information in the job cache:

#keep_jobs: 24

The number of seconds to wait when the client is requesting information

about running jobs.

#gather_job_timeout: 10

Set the default timeout for the salt command and api. The default is 5

seconds.

#timeout: 5

The loop_interval option controls the seconds for the master's maintenance
process check cycle. This process updates file server backends, cleans the
job cache and executes the scheduler.

#loop_interval: 60

Set the default outputter used by the salt command. The default is "nested".

#output: nested

To set a list of additional directories to search for salt outputters, set the
outputter_dirs option.

#outputter_dirs: []

Set the default output file used by the salt command. Default is to output
to the CLI and not to a file. Functions the same way as the "--out-file"
CLI option, only sets this to a single file for all salt commands.

#output_file: None

Return minions that timeout when running commands like test.ping

#show_timeout: True

Tell the client to display the jid when a job is published.

#show_jid: False

By default, output is colored. To disable colored output, set the color value
to False.

#color: True

Do not strip off the colored output from nested results and state outputs
(true by default).

strip_colors: False

To display a summary of the number of minions targeted, the number of
minions returned, and the number of minions that did not return, set the
cli_summary value to True. (False by default.)

#

#cli_summary: False

Set the directory used to hold unix sockets:

#sock_dir: /var/run/salt/master

The master can take a while to start up when lspci and/or dmidecode is used
to populate the grains for the master. Enable if you want to see GPU hardware
data for your master.

enable_gpu_grains: False

The master maintains a job cache. While this is a great addition, it can be
a burden on the master for larger deployments (over 5000 minions).
Disabling the job cache will make previously executed jobs unavailable to
the jobs system and is not generally recommended.

#job_cache: True

Cache minion grains, pillar and mine data via the cache subsystem in the
cachedir or a database.

#minion_data_cache: True

Cache subsystem module to use for minion data cache.

#cache: localfs

Enables a fast in-memory cache booster and sets the expiration time.

```
#memcache_expire_seconds: 0

# Set a memcache limit in items (bank + key) per cache storage (driver + driver_opts).

#memcache_max_items: 1024

# Each time a cache storage got full cleanup all the expired items not just the oldest one.

#memcache_full_cleanup: False

# Enable collecting the memcache stats and log it on `debug` log level.

#memcache_debug: False


# Store all returns in the given returner.

# Setting this option requires that any returner-specific configuration also

# be set. See various returners in salt/returners for details on required

# configuration values. (See also, event_return_queue, and event_return_queue_max_seconds below.)

#

#event_return: mysql


# On busy systems, enabling event_returns can cause a considerable load on

# the storage system for returners. Events can be queued on the master and

# stored in a batched fashion using a single transaction for multiple events.

# By default, events are not queued.

#event_return_queue: 0


# In some cases enabling event return queueing can be very helpful, but the bus

# may not be busy enough to flush the queue consistently. Setting this to a reasonable

# value (1-30 seconds) will cause the queue to be flushed when the oldest event is older

# than `event_return_queue_max_seconds` regardless of how many events are in the queue.

#event_return_queue_max_seconds: 0


# Only return events matching tags in a whitelist, supports glob matches.

#event_return_whitelist:
```

- salt/master/a_tag

- salt/run/*/ret

Store all event returns ****except**** the tags in a blacklist, supports globs.

#event_return_blacklist:

- salt/master/not_this_tag

- salt/wheel/*/ret

Passing very large events can cause the minion to consume large amounts of

memory. This value tunes the maximum size of a message allowed onto the

master event bus. The value is expressed in bytes.

#max_event_size: 1048576

Windows platforms lack posix IPC and must rely on slower TCP based inter-

process communications. Set ipc_mode to 'tcp' on such systems

#ipc_mode: ipc

Overwrite the default tcp ports used by the minion when ipc_mode is set to 'tcp'

#tcp_master_pub_port: 4510

#tcp_master_pull_port: 4511

By default, the master AES key rotates every 24 hours. The next command

following a key rotation will trigger a key refresh from the minion which may

result in minions which do not respond to the first command after a key refresh.

#

To tell the master to ping all minions immediately after an AES key refresh, set

ping_on_rotate to True. This should mitigate the issue where a minion does not

appear to initially respond after a key is rotated.

#

Note that ping_on_rotate may cause high load on the master immediately after
the key rotation event as minions reconnect. Consider this carefully if this
salt master is managing a large number of minions.

#

If disabled, it is recommended to handle this event by listening for the
'aes_key_rotate' event with the 'key' tag and acting appropriately.
ping_on_rotate: False

By default, the master deletes its cache of minion data when the key for that
minion is removed. To preserve the cache after key deletion, set
'preserve_minion_cache' to True.

#

WARNING: This may have security implications if compromised minions auth with
a previous deleted minion ID.
#preserve_minion_cache: False

Allow or deny minions from requesting their own key revocation
#allow_minion_key_revoke: True

If max_minions is used in large installations, the master might experience
high-load situations because of having to check the number of connected
minions for every authentication. This cache provides the minion-ids of
all connected minions to all MWorker-processes and greatly improves the
performance of max_minions.
con_cache: False

The master can include configuration from other files. To enable this,
pass a list of paths to this option. The paths can be either relative or
absolute; if relative, they are considered to be relative to the directory

the main master configuration file lives in (this file). Paths can make use
of shell-style globbing. If no files are matched by a path passed to this
option, then the master will log a warning message.

#

Include a config file from some other path:

include: /etc/salt/extra_config

#

Include config from several files and directories:

include:

- /etc/salt/extra_config

Large-scale tuning settings

#####

Max open files

#

Each minion connecting to the master uses AT LEAST one file descriptor, the
master subscription connection. If enough minions connect you might start
seeing on the console (and then salt-master crashes):

Too many open files (tcp_listener.cpp:335)

Aborted (core dumped)

#

By default this value will be the one of `ulimit -Hn`, ie, the hard limit for
max open files.

#

If you wish to set a different value than the default one, uncomment and
configure this setting. Remember that this value CANNOT be higher than the
hard limit. Raising the hard limit depends on your OS and/or distribution,
a good way to find the limit is to search the internet. For example:

raise max open files hard limit debian

#

#max_open_files: 100000

The number of worker threads to start. These threads are used to manage

return calls made from minions to the master. If the master seems to be

running slowly, increase the number of threads. This setting can not be

set lower than 3.

#worker_threads: 5

Set the ZeroMQ high water marks

<http://api.zeromq.org/3-2:zmq-setsockopt>

The listen queue size / backlog

#zmq_backlog: 1000

The publisher interface ZeroMQPubServerChannel

#pub_hwm: 1000

The master may allocate memory per-event and not

reclaim it.

To set a high-water mark for memory allocation, use

ipc_write_buffer to set a high-water mark for message

buffering.

Value: In bytes. Set to 'dynamic' to have Salt select

a value for you. Default is disabled.

ipc_write_buffer: 'dynamic'

These two batch settings, batch_safe_limit and batch_safe_size, are used to

automatically switch to a batch mode execution. If a command would have been
sent to more than <batch_safe_limit> minions, then run the command in
batches of <batch_safe_size>. If no batch_safe_size is specified, a default
of 8 will be used. If no batch_safe_limit is specified, then no automatic
batching will occur.
#batch_safe_limit: 100
#batch_safe_size: 8

Master stats enables stats events to be fired from the master at close
to the defined interval
#master_stats: False
#master_stats_event_iter: 60

Security settings

#####

Enable passphrase protection of Master private key. Although a string value
is acceptable; passwords should be stored in an external vaulting mechanism
and retrieved via sdb. See <https://docs.saltproject.io/en/latest/topics/sdb/>.
Passphrase protection is off by default but an example of an sdb profile and
query is as follows.

masterkeyring:

driver: keyring

service: system

#

key_pass: sdb://masterkeyring/key_pass

Enable passphrase protection of the Master signing_key. This only applies if
master_sign_pubkey is set to True. This is disabled by default.

master_sign_pubkey: True

signing_key_pass: sdb://masterkeyring/signing_pass

Enable "open mode", this mode still maintains encryption, but turns off

authentication, this is only intended for highly secure environments or for

the situation where your keys end up in a bad state. If you run in open mode

you do so at your own risk!

#open_mode: False

Enable auto_accept, this setting will automatically accept all incoming

public keys from the minions. Note that this is insecure.

#auto_accept: False

The size of key that should be generated when creating new keys.

#keysize: 2048

Time in minutes that an incoming public key with a matching name found in

pki_dir/minion_autosign/keyid is automatically accepted. Expired autosign keys

are removed when the master checks the minion_autosign directory.

0 equals no timeout

autosign_timeout: 120

If the autosign_file is specified, incoming keys specified in the

autosign_file will be automatically accepted. This is insecure. Regular

expressions as well as globbing lines are supported. The file must be readonly

except for the owner. Use permissive_pki_access to allow the group write access.

#autosign_file: /etc/salt/autosign.conf

Works like autosign_file, but instead allows you to specify minion IDs for

```
# which keys will automatically be rejected. Will override both membership in
# the autosign_file and the auto_accept setting.
#autoreject_file: /etc/salt/autoreject.conf
```

```
# If the autosign_grains_dir is specified, incoming keys from minions with grain
# values matching those defined in files in this directory will be accepted
# automatically. This is insecure. Minions need to be configured to send the grains.
#autosign_grains_dir: /etc/salt/autosign_grains
```

```
# Enable permissive access to the salt keys. This allows you to run the
# master or minion as root, but have a non-root group be given access to
# your pki_dir. To make the access explicit, root must belong to the group
# you've given access to. This is potentially quite insecure. If an autosign_file
# is specified, enabling permissive_pki_access will allow group access to that
# specific file.
#permissive_pki_access: False
```

```
# Allow users on the master access to execute specific commands on minions.
# This setting should be treated with care since it opens up execution
# capabilities to non root users. By default this capability is completely
# disabled.
```

```
#publisher_acl:
```

```
# larry:
```

```
# - test.ping
```

```
# - network.*
```

```
#
```

```
# Blacklist any of the following users or modules
```

```
#
```

```
# This example would blacklist all non sudo users, including root from
```

```
# running any commands. It would also blacklist any use of the "cmd"
# module. This is completely disabled by default.
#
#
# Check the list of configured users in client ACL against users on the
# system and throw errors if they do not exist.
#client_acl_verify: True
#
#publisher_acl_blacklist:
# users:
# - root
# - '^(?!sudo_).*'$ # all non sudo users
# modules:
# - cmd

# Enforce publisher_acl & publisher_acl_blacklist when users have sudo
# access to the salt command.
#
#sudo_acl: False

# The external auth system uses the Salt auth modules to authenticate and
# validate users to access areas of the Salt system.
#external_auth:
# pam:
# fred:
# - test.*
#
# Time (in seconds) for a newly generated token to live. Default: 12 hours
#token_expire: 43200
```

```
#
# Allow eauth users to specify the expiry time of the tokens they generate.
# A boolean applies to all users or a dictionary of whitelisted eauth backends
# and usernames may be given.
# token_expire_user_override:
#   pam:
#     - fred
#     - tom
#   ldap:
#     - gary
#
#token_expire_user_override: False

# Set to True to enable keeping the calculated user's auth list in the token
# file. This is disabled by default and the auth list is calculated or requested
# from the eauth driver each time.
#keep_acl_in_token: False

# Auth subsystem module to use to get authorized access list for a user. By default it's
# the same module used for external authentication.
#eauth_acl_module: django

# Allow minions to push files to the master. This is disabled by default, for
# security purposes.
#file_recv: False

# Set a hard-limit on the size of the files that can be pushed to the master.
# It will be interpreted as megabytes. Default: 100
#file_recv_max_size: 100
```

```
# Signature verification on messages published from the master.
# This causes the master to cryptographically sign all messages published to its event
# bus, and minions then verify that signature before acting on the message.
#
# This is False by default.
#
# Note that to facilitate interoperability with masters and minions that are different
# versions, if sign_pub_messages is True but a message is received by a minion with
# no signature, it will still be accepted, and a warning message will be logged.
# Conversely, if sign_pub_messages is False, but a minion receives a signed
# message it will be accepted, the signature will not be checked, and a warning message
# will be logged. This behavior went away in Salt 2014.1.0 and these two situations
# will cause minion to throw an exception and drop the message.
# sign_pub_messages: False

# Signature verification on messages published from minions
# This requires that minions cryptographically sign the messages they
# publish to the master. If minions are not signing, then log this information
# at loglevel 'INFO' and drop the message without acting on it.
# require_minion_sign_messages: False

# The below will drop messages when their signatures do not validate.
# Note that when this option is False but `require_minion_sign_messages` is True
# minions MUST sign their messages but the validity of their signatures
# is ignored.
# These two config options exist so a Salt infrastructure can be moved
# to signing minion messages gradually.
# drop_messages_signature_fail: False
```



```
# Use TLS/SSL encrypted connection between master and minion.

# Can be set to a dictionary containing keyword arguments corresponding to Python's
# 'ssl.wrap_socket' method.

# Default is None.

#ssl:

#  keyfile: <path_to_keyfile>
#  certfile: <path_to_certfile>
#  ssl_version: PROTOCOL_TLSv1_2


#####  Salt-SSH Configuration  #####
#####

# Define the default salt-ssh roster module to use

#roster: flat


# Pass in an alternative location for the salt-ssh `flat` roster file

#roster_file: /etc/salt/roster


# Define locations for `flat` roster files so they can be chosen when using Salt API.
# An administrator can place roster files into these locations. Then when
# calling Salt API, parameter 'roster_file' should contain a relative path to
# these locations. That is, "roster_file=/foo/roster" will be resolved as
# "/etc/salt/roster.d/foo/roster" etc. This feature prevents passing insecure
# custom rosters through the Salt API.

#

#rosters:

# - /etc/salt/roster.d
# - /opt/salt/some/more/rosters
```

The ssh password to log in with.

#ssh_passwd: "

#The target system's ssh port number.

#ssh_port: 22

Comma-separated list of ports to scan.

#ssh_scan_ports: 22

Scanning socket timeout for salt-ssh.

#ssh_scan_timeout: 0.01

Boolean to run command via sudo.

#ssh_sudo: False

Boolean to run ssh_pre_flight script defined in roster. By default

the script will only run if the thin_dir does not exist on the targeted

minion. This forces the script to run regardless of the thin dir existing

or not.

#ssh_run_pre_flight: True

Number of seconds to wait for a response when establishing an SSH connection.

#ssh_timeout: 60

The user to log in as.

#ssh_user: root

The log file of the salt-ssh command:

#ssh_log_file: /var/log/salt/ssh

Pass in minion option overrides that will be inserted into the SHIM for

salt-ssh calls. The local minion config is not used for salt-ssh. Can be

overridden on a per-minion basis in the roster (`minion_opts`)

#ssh_minion_opts:

gpg_keydir: /root/gpg

Set this to True to default to using ~/.ssh/id_rsa for salt-ssh

authentication with minions

#ssh_use_home_key: False

Set this to True to default salt-ssh to run with ``-o IdentitiesOnly=yes``.

This option is intended for situations where the ssh-agent offers many

different identities and allows ssh to ignore those identities and use the

only one specified in options.

#ssh_identities_only: False

List-only nodegroups for salt-ssh. Each group must be formed as either a

comma-separated list, or a YAML list. This option is useful to group minions

into easy-to-target groups when using salt-ssh. These groups can then be

targeted with the normal -N argument to salt-ssh.

#ssh_list_nodegroups: {}

salt-ssh has the ability to update the flat roster file if a minion is not

found in the roster. Set this to True to enable it.

#ssh_update_roster: False

Master Module Management

#####

Manage how master side modules are loaded.

Add any additional locations to look for master runners:

#runner_dirs: []

Add any additional locations to look for master utils:

#utils_dirs: []

Enable Cython for master side modules:

#cython_enable: False

State System settings

#####

The state system uses a "top" file to tell the minions what environment to

use and what modules to use. The state_top file is defined relative to the

root of the base environment as defined in "File Server settings" below.

#state_top: top.sls

The master_tops option replaces the external_nodes option by creating

a pluggable system for the generation of external top data. The external_nodes

option is deprecated by the master_tops option.

#

To gain the capabilities of the classic external_nodes system, use the

following configuration:

master_tops:

ext_nodes: <Shell command which returns yaml>

#

#master_tops: {}

The renderer to use on the minions to render the state data

#renderer: jinja|yaml

Default Jinja environment options for all templates except sls templates

#jinja_env:

block_start_string: '{%'

block_end_string: '%}'

variable_start_string: '{{'

variable_end_string: '}}'

comment_start_string: '{#'

comment_end_string: '#}'

line_statement_prefix:

line_comment_prefix:

trim_blocks: False

lstrip_blocks: False

newline_sequence: '\n'

keep_trailing_newline: False

Jinja environment options for sls templates

#jinja_sls_env:

block_start_string: '{%'

block_end_string: '%}'

variable_start_string: '{{'

variable_end_string: '}}'

comment_start_string: '{#'

comment_end_string: '#}'

line_statement_prefix:

line_comment_prefix:

```
# trim_blocks: False
# lstrip_blocks: False
# newline_sequence: '\n'
# keep_trailing_newline: False
```

```
# The failhard option tells the minions to stop immediately after the first
# failure detected in the state execution, defaults to False
#failhard: False
```

```
# The state_verbose and state_output settings can be used to change the way
# state system data is printed to the display. By default all data is printed.
# The state_verbose setting can be set to True or False, when set to False
# all data that has a result of True and no changes will be suppressed.
#state_verbose: True
```

```
# The state_output setting controls which results will be output full multi line
# full, terse - each state will be full/terse
# mixed - only states with errors will be full
# changes - states with changes and errors will be full
# full_id, mixed_id, changes_id and terse_id are also allowed;
# when set, the state ID will be used as name in the output
#state_output: full
```

```
# The state_output_diff setting changes whether or not the output from
# successful states is returned. Useful when even the terse output of these
# states is cluttering the logs. Set it to True to ignore them.
#state_output_diff: False
```

```
# The state_output_profile setting changes whether profile information
```

will be shown for each state run.

#state_output_profile: True

Automatically aggregate all states that have support for mod_aggregate by

setting to 'True'. Or pass a list of state module names to automatically

aggregate just those types.

#

state_aggregate:

- pkg

#

#state_aggregate: False

Send progress events as each function in a state run completes execution

by setting to 'True'. Progress events are in the format

'salt/job/<JID>/prog/<MID>/<RUN NUM>'.

#state_events: False

File Server settings

#####

Salt runs a lightweight file server written in zeromq to deliver files to

minions. This file server is built into the master daemon and does not

require a dedicated port.

The file server works on environments passed to the master, each environment

can have multiple root directories, the subdirectories in the multiple file

roots cannot match, otherwise the downloaded files will not be able to be

reliably ensured. A base environment is required to house the top file.

Example:

file_roots:

```
# base:
# - /srv/salt/

# dev:
# - /srv/salt/dev/services
# - /srv/salt/dev/states

# prod:
# - /srv/salt/prod/services
# - /srv/salt/prod/states
#

#file_roots:
# base:
# - /srv/salt
#

# The master_roots setting configures a master-only copy of the file_roots dictionary,
# used by the state compiler.

#master_roots:
# base:
# - /srv/salt-master

# When using multiple environments, each with their own top file, the
# default behaviour is an unordered merge. To prevent top files from
# being merged together and instead to only use the top file from the
# requested environment, set this value to 'same'.

#top_file_merging_strategy: merge

# To specify the order in which environments are merged, set the ordering
# in the env_order option. Given a conflict, the last matching value will
# win.
```



```
#env_order: ['base', 'dev', 'prod']
```

```
# If top_file_merging_strategy is set to 'same' and an environment does not  
# contain a top file, the top file in the environment specified by default_top  
# will be used instead.
```

```
#default_top: base
```

```
# The hash_type is the hash to use when discovering the hash of a file on  
# the master server. The default is sha256, but md5, sha1, sha224, sha384 and  
# sha512 are also supported.
```

```
#
```

```
# WARNING: While md5 and sha1 are also supported, do not use them due to the  
# high chance of possible collisions and thus security breach.
```

```
#
```

```
# Prior to changing this value, the master should be stopped and all Salt  
# caches should be cleared.
```

```
#hash_type: sha256
```

```
# The buffer size in the file server can be adjusted here:
```

```
#file_buffer_size: 1048576
```

```
# A regular expression (or a list of expressions) that will be matched  
# against the file path before syncing the modules and states to the minions.  
# This includes files affected by the file.recurse state.
```

```
# For example, if you manage your custom modules and states in subversion  
# and don't want all the '.svn' folders and content synced to your minions,  
# you could set this to '\.svn($|/)'. By default nothing is ignored.
```

```
#file_ignore_regex:
```

```
# - '\.svn($|/)'
```

```
# - '/\.git($|/).'
```

```
# A file glob (or list of file globs) that will be matched against the file  
# path before syncing the modules and states to the minions. This is similar  
# to file_ignore_regex above, but works on globs instead of regex. By default  
# nothing is ignored.
```

```
# file_ignore_glob:
```

```
# - '*.pyc'
```

```
# - '*/somefolder/*.bak'
```

```
# - '*.swp'
```

```
# File Server Backend
```

```
#
```

```
# Salt supports a modular fileserver backend system, this system allows  
# the salt master to link directly to third party systems to gather and  
# manage the files available to minions. Multiple backends can be  
# configured and will be searched for the requested file in the order in which  
# they are defined here. The default setting only enables the standard backend  
# "roots" which uses the "file_roots" option.
```

```
#fileserver_backend:
```

```
# - roots
```

```
#
```

```
# To use multiple backends list them in the order they are searched:
```

```
#fileserver_backend:
```

```
# - git
```

```
# - roots
```

```
#
```

```
# Uncomment the line below if you do not want the file_server to follow
```

```
# symlinks when walking the filesystem tree. This is set to True
```

```
# by default. Currently this only applies to the default roots

# fileserver_backend.

#fileserver_followsymlinks: False

#

# Uncomment the line below if you do not want symlinks to be
# treated as the files they are pointing to. By default this is set to
# False. By uncommenting the line below, any detected symlink while listing
# files on the Master will not be returned to the Minion.

#fileserver_ignoresymlinks: True

#

# By default, the Salt fileserver recurses fully into all defined environments
# to attempt to find files. To limit this behavior so that the fileserver only
# traverses directories with SLS files and special Salt directories like _modules,
# enable the option below. This might be useful for installations where a file root
# has a very large number of files and performance is impacted. Default is False.

# fileserver_limit_traversal: False

#

# The fileserver can fire events off every time the fileserver is updated,
# these are disabled by default, but can be easily turned on by setting this
# flag to True

#fileserver_events: False


# Git File Server Backend Configuration

#

# Optional parameter used to specify the provider to be used for gitfs. Must be
# either pygit2 or gitpython. If unset, then both will be tried (in that
# order), and the first one with a compatible version installed will be the
# provider that is used.

#
```

#gitfs_provider: pygit2

Along with gitfs_password, is used to authenticate to HTTPS remotes.

gitfs_user: "

Along with gitfs_user, is used to authenticate to HTTPS remotes.

This parameter is not required if the repository does not use authentication.

#gitfs_password: "

By default, Salt will not authenticate to an HTTP (non-HTTPS) remote.

This parameter enables authentication over HTTP. Enable this at your own risk.

#gitfs_insecure_auth: False

Along with gitfs_privkey (and optionally gitfs_passphrase), is used to

authenticate to SSH remotes. This parameter (or its per-remote counterpart)

is required for SSH remotes.

#gitfs_pubkey: "

Along with gitfs_pubkey (and optionally gitfs_passphrase), is used to

authenticate to SSH remotes. This parameter (or its per-remote counterpart)

is required for SSH remotes.

#gitfs_privkey: "

This parameter is optional, required only when the SSH key being used to

authenticate is protected by a passphrase.

#gitfs_passphrase: "

When using the git fileserver backend at least one git remote needs to be

defined. The user running the salt master will need read access to the repo.

```
#

# The repos will be searched in order to find the file requested by a client
# and the first repo to have the file will return it.

# When using the git backend branches and tags are translated into salt
# environments.

# Note: file:// repos will be treated as a remote, so refs you want used must
# exist in that repo as *local* refs.

#gitfs_remotes:
# - git://github.com/saltstack/salt-states.git
# - file:///var/git/saltmaster

#

# The gitfs_ssl_verify option specifies whether to ignore ssl certificate
# errors when contacting the gitfs backend. You might want to set this to
# false if you're using a git backend that uses a self-signed certificate but
# keep in mind that setting this flag to anything other than the default of True
# is a security concern, you may want to try using the ssh transport.

#gitfs_ssl_verify: True

#

# The gitfs_root option gives the ability to serve files from a subdirectory
# within the repository. The path is defined relative to the root of the
# repository and defaults to the repository root.

#gitfs_root: somefolder/otherfolder

#

# The refsspecs fetched by gitfs remotes

#gitfs_refsspecs:
# - '+refs/heads/*:refs/remotes/origin/*'
# - '+refs/tags/*:refs/tags/*'

#

#
```

```
##### Pillar settings #####
```

```
#####
```

```
# Salt Pillars allow for the building of global data that can be made selectively
# available to different minions based on minion grain filtering. The Salt
# Pillar is laid out in the same fashion as the file server, with environments,
# a top file and sls files. However, pillar data does not need to be in the
# highstate format, and is generally just key/value pairs.
```

```
#pillar_roots:
```

```
# base:
```

```
# - /srv/pillar
```

```
#
```

```
#ext_pillar:
```

```
# - hiera: /etc/hiera.yaml
```

```
# - cmd_yaml: cat /etc/salt/yaml
```

```
# A list of paths to be recursively decrypted during pillar compilation.
```

```
# Entries in this list can be formatted either as a simple string, or as a
```

```
# key/value pair, with the key being the pillar location, and the value being
```

```
# the renderer to use for pillar decryption. If the former is used, the
```

```
# renderer specified by decrypt_pillar_default will be used.
```

```
#decrypt_pillar:
```

```
# - 'foo:bar': gpg
```

```
# - 'lorem:ipsum:dolor'
```

```
# The delimiter used to distinguish nested data structures in the
```

```
# decrypt_pillar option.
```

```
#decrypt_pillar_delimiter: ':'
```

```
# The default renderer used for decryption, if one is not specified for a given
# pillar key in decrypt_pillar.
#decrypt_pillar_default: gpg
```

```
# List of renderers which are permitted to be used for pillar decryption.
#decrypt_pillar_renderers:
# - gpg
```

```
# The ext_pillar_first option allows for external pillar sources to populate
# before file system pillar. This allows for targeting file system pillar from
# ext_pillar.
#ext_pillar_first: False
```

```
# The external pillars permitted to be used on-demand using pillar.ext
#on_demand_ext_pillar:
# - libvirt
# - virtkey
```

```
# The pillar_gitfs_ssl_verify option specifies whether to ignore ssl certificate
# errors when contacting the pillar gitfs backend. You might want to set this to
# false if you're using a git backend that uses a self-signed certificate but
# keep in mind that setting this flag to anything other than the default of True
# is a security concern, you may want to try using the ssh transport.
#pillar_gitfs_ssl_verify: True
```

```
# The pillar_opts option adds the master configuration file data to a dict in
# the pillar called "master". This is used to set simple configurations in the
# master config file that can then be used on minions.
#pillar_opts: False
```

The pillar_safe_render_error option prevents the master from passing pillar
render errors to the minion. This is set on by default because the error could
contain templating data which would give that minion information it shouldn't
have, like a password! When set true the error message will only show:
Rendering SLS 'my.sls' failed. Please see master log for details.
#pillar_safe_render_error: True

The pillar_source_merging_strategy option allows you to configure merging strategy
between different sources. It accepts five values: none, recurse, aggregate, overwrite,
or smart. None will not do any merging at all. Recurse will merge recursively mapping of data.
Aggregate instructs aggregation of elements between sources that use the #!yamlex renderer.
Overwrite
will overwrite elements according the order in which they are processed. This is
behavior of the 2014.1 branch and earlier. Smart guesses the best strategy based
on the "renderer" setting and is the default value.
#pillar_source_merging_strategy: smart

Recursively merge lists by aggregating them instead of replacing them.
#pillar_merge_lists: False

Set this option to True to force the pillarenv to be the same as the effective
saltenv when running states. If pillarenv is specified this option will be
ignored.
#pillarenv_from_saltenv: False

Set this option to 'True' to force a 'KeyError' to be raised whenever an
attempt to retrieve a named value from pillar fails. When this option is set
to 'False', the failed attempt returns an empty string. Default is 'False'.

#pillar_raise_on_missing: False

Git External Pillar (git_pillar) Configuration Options

#

Specify the provider to be used for git_pillar. Must be either pygit2 or
gitpython. If unset, then both will be tried in that same order, and the
first one with a compatible version installed will be the provider that
is used.

#git_pillar_provider: pygit2

If the desired branch matches this value, and the environment is omitted
from the git_pillar configuration, then the environment for that git_pillar
remote will be base.

#git_pillar_base: master

If the branch is omitted from a git_pillar remote, then this branch will
be used instead

#git_pillar_branch: master

Environment to use for git_pillar remotes. This is normally derived from
the branch/tag (or from a per-remote env parameter), but if set this will
override the process of deriving the env from the branch/tag name.

#git_pillar_env: "

Path relative to the root of the repository where the git_pillar top file
and SLS files are located.

#git_pillar_root: "

Specifies whether or not to ignore SSL certificate errors when contacting

the remote repository.

#git_pillar_ssl_verify: False

When set to False, if there is an update/checkout lock for a git_pillar

remote and the pid written to it is not running on the master, the lock

file will be automatically cleared and a new lock will be obtained.

#git_pillar_global_lock: True

Git External Pillar Authentication Options

#

Along with git_pillar_password, is used to authenticate to HTTPS remotes.

#git_pillar_user: "

Along with git_pillar_user, is used to authenticate to HTTPS remotes.

This parameter is not required if the repository does not use authentication.

#git_pillar_password: "

By default, Salt will not authenticate to an HTTP (non-HTTPS) remote.

This parameter enables authentication over HTTP.

#git_pillar_insecure_auth: False

Along with git_pillar_privkey (and optionally git_pillar_passphrase),

is used to authenticate to SSH remotes.

#git_pillar_pubkey: "

Along with git_pillar_pubkey (and optionally git_pillar_passphrase),

is used to authenticate to SSH remotes.

#git_pillar_privkey: "

This parameter is optional, required only when the SSH key being used

to authenticate is protected by a passphrase.

#git_pillar_passphrase: "

The refsspecs fetched by git_pillar remotes

#git_pillar_refsspecs:

- '+refs/heads/*:refs/remotes/origin/*'

- '+refs/tags/*:refs/tags/*'

A master can cache pillars locally to bypass the expense of having to render them

for each minion on every request. This feature should only be enabled in cases

where pillar rendering time is known to be unsatisfactory and any attendant security

concerns about storing pillars in a master cache have been addressed.

#

When enabling this feature, be certain to read through the additional ``pillar_cache_*``

configuration options to fully understand the tunable parameters and their implications.

#

Note: setting ``pillar_cache: True`` has no effect on targeting Minions with Pillars.

See <https://docs.saltproject.io/en/latest/topics/targeting/pillar.html>

#pillar_cache: False

If and only if a master has set ``pillar_cache: True``, the cache TTL controls the amount

of time, in seconds, before the cache is considered invalid by a master and a fresh

pillar is recompiled and stored.

#pillar_cache_ttl: 3600

If and only if a master has set `pillar_cache: True`, one of several storage providers

can be utilized.

#

```
# `disk`: The default storage backend. This caches rendered pillars to the master cache.
#     Rendered pillars are serialized and deserialized as msgpack structures for speed.
#     Note that pillars are stored UNENCRYPTED. Ensure that the master cache
#     has permissions set appropriately. (Same defaults are provided.)
#
# memory: [EXPERIMENTAL] An optional backend for pillar caches which uses a pure-Python
#     in-memory data structure for maximal performance. There are several caveats,
#     however. First, because each master worker contains its own in-memory cache,
#     there is no guarantee of cache consistency between minion requests. This
#     works best in situations where the pillar rarely if ever changes. Secondly,
#     and perhaps more importantly, this means that unencrypted pillars will
#     be accessible to any process which can examine the memory of the ``salt-master``!
#     This may represent a substantial security risk.
#
#pillar_cache_backend: disk

# A master can also cache GPG data locally to bypass the expense of having to render them
# for each minion on every request. This feature should only be enabled in cases
# where pillar rendering time is known to be unsatisfactory and any attendant security
# concerns about storing decrypted GPG data in a master cache have been addressed.
#
# When enabling this feature, be certain to read through the additional ``gpg_cache_*``
# configuration options to fully understand the tunable parameters and their implications.
#gpg_cache: False

# If and only if a master has set ``gpg_cache: True``, the cache TTL controls the amount
# of time, in seconds, before the cache is considered invalid by a master and a fresh
# pillar is recompiled and stored.
#gpg_cache_ttl: 86400
```

```
# If and only if a master has set `gpg_cache: True`, one of several storage providers
# can be utilized. Available options are the same as ``pillar_cache_backend``.
#gpg_cache_backend: disk
```

```
##### Reactor Settings #####
```

```
#####
```

```
# Define a salt reactor. See https://docs.saltproject.io/en/latest/topics/reactor/
#reactor: []
```

```
#Set the TTL for the cache of the reactor configuration.
#reactor_refresh_interval: 60
```

```
#Configure the number of workers for the runner/wheel in the reactor.
#reactor_worker_threads: 10
```

```
#Define the queue size for workers in the reactor.
#reactor_worker_hwm: 10000
```

```
##### Syndic settings #####
```

```
#####
```

```
# The Salt syndic is used to pass commands through a master from a higher
# master. Using the syndic is simple. If this is a master that will have
# syndic servers(s) below it, then set the "order_masters" setting to True.
#
```

```
# If this is a master that will be running a syndic daemon for passthrough, then
# the "syndic_master" setting needs to be set to the location of the master server
```

to receive commands from.

Set the order_masters setting to True if this master will command lower

masters' syndic interfaces.

#order_masters: False

If this master will be running a salt syndic daemon, syndic_master tells

this master where to receive commands from.

#syndic_master: masterofmasters

This is the 'ret_port' of the MasterOfMaster:

#syndic_master_port: 4506

PID file of the syndic daemon:

#syndic_pidfile: /var/run/salt-syndic.pid

The log file of the salt-syndic daemon:

#syndic_log_file: /var/log/salt/syndic

The behaviour of the multi-syndic when connection to a master of masters failed.

Can specify ``random`` (default) or ``ordered``. If set to ``random``, masters

will be iterated in random order. If ``ordered`` is specified, the configured

order will be used.

#syndic_failover: random

The number of seconds for the salt client to wait for additional syndics to

check in with their lists of expected minions before giving up.

#syndic_wait: 5

Peer Publish settings

#####

Salt minions can send commands to other minions, but only if the minion is
allowed to. By default "Peer Publication" is disabled, and when enabled it
is enabled for specific minions and specific commands. This allows secure
compartmentalization of commands based on individual minions.

The configuration uses regular expressions to match minions and then a list
of regular expressions to match functions. The following will allow the
minion authenticated as foo.example.com to execute functions from the test
and pkg modules.

#peer:

foo.example.com:

- test.*

- pkg.*

#

This will allow all minions to execute all commands:

#peer:

.*:

- .*

#

This is not recommended, since it would allow anyone who gets root on any
single minion to instantly have root on all of the minions!

Minions can also be allowed to execute runners from the salt master.

Since executing a runner from the minion could be considered a security risk,
it needs to be enabled. This setting functions just like the peer setting
except that it opens up runners instead of module functions.

```

#
# All peer runner support is turned off by default and must be enabled before
# using. This will enable all peer runners for all minions:
#peer_run:
# .*:
# - .*
#
# To enable just the manage.up runner for the minion foo.example.com:
#peer_run:
# foo.example.com:
# - manage.up
#
#
#####      Mine settings      #####
#####

# Restrict mine.get access from minions. By default any minion has a full access
# to get all mine data from master cache. In acl definition below, only pcre matches
# are allowed.
# mine_get:
# .*:
# - .*
#
# The example below enables minion foo.example.com to get 'network.interfaces' mine
# data only, minions web* to get all network.* and disk.* mine data and all other
# minions won't get any mine data.
# mine_get:
# foo.example.com:
# - network.interfaces
# web.*:

```


- network.*

- disk.*

Logging settings

#####

The location of the master log file

The master log can be sent to a regular file, local path name, or network

location. Remote logging works best when configured to use rsyslogd(8) (e.g.:

``file:///dev/log``), with rsyslogd(8) configured for network logging. The URI

format is: <file|udp|tcp>://<host|socketpath>:<port-if-required>/<log-facility>

#log_file: /var/log/salt/master

#log_file: file:///dev/log

#log_file: udp://loghost:10514

#log_file: /var/log/salt/master

#key_logfile: /var/log/salt/key

The level of messages to send to the console.

One of 'garbage', 'trace', 'debug', 'info', 'warning', 'error', 'critical'.

#

The following log levels are considered INSECURE and may log sensitive data:

['garbage', 'trace', 'debug']

#

#log_level: warning

The level of messages to send to the log file.

One of 'garbage', 'trace', 'debug', 'info', 'warning', 'error', 'critical'.

If using 'log_granular_levels' this must be set to the highest desired level.

#log_level_logfile: warning

The date and time format used in log messages. Allowed date/time formatting

can be seen here: <http://docs.python.org/library/time.html#time.strftime>

#log_datefmt: '%H:%M:%S'

#log_datefmt_logfile: '%Y-%m-%d %H:%M:%S'

The format of the console logging messages. Allowed formatting options can

be seen here: <http://docs.python.org/library/logging.html#logrecord-attributes>

#

Console log colors are specified by these additional formatters:

#

%(colorlevel)s

%(colorname)s

%(colorprocess)s

%(colormsg)s

#

Since it is desirable to include the surrounding brackets, '[' and ']', in

the coloring of the messages, these color formatters also include padding as

well. Color LogRecord attributes are only available for console logging.

#

#log_fmt_console: '%(colorlevel)s %(colormsg)s'

#log_fmt_console: '[(levelname)-8s] %(message)s'

#

#log_fmt_logfile: '%(asctime)s,%(msecs)03d [%(name)-17s]%(levelname)-8s] %(message)s'

This can be used to control logging levels more specifically. This

example sets the main salt library at the 'warning' level, but sets

'salt.modules' to log at the 'debug' level:

```
# log_granular_levels:
# 'salt': 'warning'
# 'salt.modules': 'debug'
#
#log_granular_levels: {}
```

Node Groups

#####

```
# Node groups allow for logical groupings of minion nodes. A group consists of
# a group name and a compound target. Nodgroups can reference other nodegroups
# with 'N@' classifier. Ensure that you do not have circular references.
#
```

```
#nodegroups:
```

```
# group1: 'L@foo.domain.com,bar.domain.com,baz.domain.com or bl*.domain.com'
# group2: 'G@os:Debian and foo.domain.com'
# group3: 'G@os:Debian and N@group1'
# group4:
# - 'G@foo:bar'
# - 'or'
# - 'G@foo:baz'
```

Range Cluster settings

#####

```
# The range server (and optional port) that serves your cluster information
# https://github.com/ytoolshed/range/wiki/%22yamlfile%22-module-file-spec
#
#range_server: range:80
```

Windows Software Repo settings

#####

Location of the repo on the master:

#winrepo_dir_ng: '/srv/salt/win/repo-ng'

#

List of git repositories to include with the local repo:

#winrepo_remotes_ng:

- 'https://github.com/saltstack/salt-winrepo-ng.git'

Windows Software Repo settings - Pre 2015.8

#####

Legacy repo settings for pre-2015.8 Windows minions.

#

Location of the repo on the master:

#winrepo_dir: '/srv/salt/win/repo'

#

Location of the master's repo cache file:

#winrepo_mastercachefile: '/srv/salt/win/repo/winrepo.p'

#

List of git repositories to include with the local repo:

#winrepo_remotes:

- 'https://github.com/saltstack/salt-winrepo.git'

The refsspecs fetched by winrepo remotes

#winrepo_refsspecs:

- '+refs/heads/*:refs/remotes/origin/*'

```
# - '+refs/tags/*:refs/tags/*'
```

```
#
```

```
##### Returner settings #####
```

```
#####
```

```
# Which returner(s) will be used for minion's result:
```

```
#return: mysql
```

```
##### Miscellaneous settings #####
```

```
#####
```

```
# Default match type for filtering events tags: startswith, endswith, find, regex, fnmatch
```

```
#event_match_type: startswith
```

```
# Save runner returns to the job cache
```

```
#runner_returns: True
```

```
# Permanently include any available Python 3rd party modules into thin and minimal Salt
```

```
# when they are generated for Salt-SSH or other purposes.
```

```
# The modules should be named by the names they are actually imported inside the Python.
```

```
# The value of the parameters can be either one module or a comma separated list of them.
```

```
#thin_extra_mods: foo,bar
```

```
#min_extra_mods: foo,bar,baz
```

```
##### Keepalive settings #####
```

```
#####
```

```
# Warning: Failure to set TCP keepalives on the salt-master can result in
```

```
# not detecting the loss of a minion when the connection is lost or when
```

its host has been terminated without first closing the socket.

Salt's Presence System depends on this connection status to know if a minion
is "present".

ZeroMQ now includes support for configuring SO_KEEPALIVE if supported by
the OS. If connections between the minion and the master pass through
a state tracking device such as a firewall or VPN gateway, there is
the risk that it could tear down the connection the master and minion
without informing either party that their connection has been taken away.
Enabling TCP Keepalives prevents this from happening.

Overall state of TCP Keepalives, enable (1 or True), disable (0 or False)
or leave to the OS defaults (-1), on Linux, typically disabled. Default True, enabled.
#tcp_keepalive: True

How long before the first keepalive should be sent in seconds. Default 300
to send the first keepalive after 5 minutes, OS default (-1) is typically 7200 seconds
on Linux see /proc/sys/net/ipv4/tcp_keepalive_time.
#tcp_keepalive_idle: 300

How many lost probes are needed to consider the connection lost. Default -1
to use OS defaults, typically 9 on Linux, see /proc/sys/net/ipv4/tcp_keepalive_probes.
#tcp_keepalive_cnt: -1

How often, in seconds, to send keepalives after the first one. Default -1 to
use OS defaults, typically 75 seconds on Linux, see
/proc/sys/net/ipv4/tcp_keepalive_intvl.
#tcp_keepalive_intvl: -1

```

#####      NetAPI settings      #####

#####

# Allow the raw_shell parameter to be used when calling Salt SSH client via API

#netapi_allow_raw_shell: True

EXAMPLE MINION CONFIGURATION FILE

##### Primary configuration settings #####

#####

# This configuration file is used to manage the behavior of the Salt Minion.

# With the exception of the location of the Salt Master Server, values that are

# commented out but have an empty line after the comment are defaults that need

# not be set in the config. If there is no blank line after the comment, the

# value is presented as an example and is not the default.


# Per default the minion will automatically include all config files

# from minion.d/*.conf (minion.d is a directory in the same directory

# as the main minion config file).

#default_include: minion.d/*.conf


# Set the location of the salt master server. If the master server cannot be

# resolved, then the minion will fail to start.

#master: salt


# Set http proxy information for the minion when doing requests

#proxy_host:

#proxy_port:

#proxy_username:

#proxy_password:


# List of hosts to bypass HTTP proxy. This key does nothing unless proxy_host etc is

```

configured, it does not support any kind of wildcards.

#no_proxy: []

If multiple masters are specified in the 'master' setting, the default behavior
is to always try to connect to them in the order they are listed. If random_master
is set to True, the order will be randomized upon Minion startup instead. This can
be helpful in distributing the load of many minions executing salt-call requests,
for example, from a cron job. If only one master is listed, this setting is ignored
and a warning will be logged.

#random_master: False

NOTE: Deprecated in Salt 2019.2.0. Use 'random_master' instead.

#master_shuffle: False

Minions can connect to multiple masters simultaneously (all masters
are "hot"), or can be configured to failover if a master becomes
unavailable. Multiple hot masters are configured by setting this
value to "str". Failover masters can be requested by setting
to "failover". MAKE SURE TO SET master_alive_interval if you are
using failover.

Setting master_type to 'disable' lets you have a running minion (with engines and
beacons) without a master connection

master_type: str

Poll interval in seconds for checking if the master is still there. Only
respected if master_type above is "failover". To disable the interval entirely,
set the value to -1. (This may be necessary on machines which have high numbers
of TCP connections, such as load balancers.)

master_alive_interval: 30

If the minion is in multi-master mode and the master_type configuration option
is set to "failover", this setting can be set to "True" to force the minion
to fail back to the first master in the list if the first master is back online.
#master_failback: False

If the minion is in multi-master mode, the "master_type" configuration is set to
"failover", and the "master_failback" option is enabled, the master failback
interval can be set to ping the top master with this interval, in seconds.
#master_failback_interval: 0

Set whether the minion should connect to the master via IPv6:
#ipv6: False

Set the number of seconds to wait before attempting to resolve
the master hostname if name resolution fails. Defaults to 30 seconds.
Set to zero if the minion should shutdown and not retry.
retry_dns: 30

Set the number of times to attempt to resolve
the master hostname if name resolution fails. Defaults to None,
which will attempt the resolution indefinitely.
retry_dns_count: 3

Set the port used by the master reply and authentication server.
#master_port: 4506

The user to run salt.
#user: root

The user to run salt remote execution commands as via sudo. If this option is
enabled then sudo will be used to change the active user executing the remote
command. If enabled the user will need to be allowed access via the sudoers
file for the user that the salt minion is configured to run as. The most
common option would be to use the root user. If this option is set the user
option should also be set to a non-root user. If migrating from a root minion
to a non root minion the minion cache should be cleared and the minion pki
directory will need to be changed to the ownership of the new user.
#sudo_user: root

Specify the location of the daemon process ID file.
#pidfile: /var/run/salt-minion.pid

The root directory prepended to these options: pki_dir, cachedir, log_file,
sock_dir, pidfile.
#root_dir: /

The path to the minion's configuration file.
#conf_file: /etc/salt/minion

The directory to store the pki information in
#pki_dir: /etc/salt/pki/minion

Explicitly declare the id for this minion to use, if left commented the id
will be the hostname as returned by the python call: socket.getfqdn()
Since salt uses detached ids it is possible to run multiple minions on the
same machine but with different ids, this can be useful for salt compute
clusters.

#id:

Cache the minion id to a file when the minion's id is not statically defined
in the minion config. Defaults to "True". This setting prevents potential
problems when automatic minion id resolution changes, which can cause the
minion to lose connection with the master. To turn off minion id caching,
set this config to ``False``.

#minion_id_caching: True

Convert minion id to lowercase when it is being generated. Helpful when some
hosts get the minion id in uppercase. Cached ids will remain the same and
not converted. For example, Windows minions often have uppercase minion
names when they are set up but not always. To turn on, set this config to
``True``.

#minion_id_lowercase: False

Append a domain to a hostname in the event that it does not exist. This is
useful for systems where socket.getfqdn() does not actually result in a
FQDN (for instance, Solaris).

#append_domain:

Custom static grains for this minion can be specified here and used in SLS
files just like all other grains. This example sets 4 custom grains, with
the 'roles' grain having two values that can be matched against.

#grains:

roles:

- webserver

- memcache

deployment: datacenter4

```
# cabinet: 13
# cab_u: 14-15
#
# Where cache data goes.
# This data may contain sensitive data and should be protected accordingly.
#cachedir: /var/cache/salt/minion

# Append minion_id to these directories. Helps with
# multiple proxies and minions running on the same machine.
# Allowed elements in the list: pki_dir, cachedir, extension_modules
# Normally not needed unless running several proxies and/or minions on the same machine
# Defaults to ['cachedir'] for proxies, [] (empty list) for regular minions
#append_minionid_config_dirs:

# Verify and set permissions on configuration directories at startup.
#verify_env: True

# The minion can locally cache the return data from jobs sent to it, this
# can be a good way to keep track of jobs the minion has executed
# (on the minion side). By default this feature is disabled, to enable, set
# cache_jobs to True.
#cache_jobs: False

# Set the directory used to hold unix sockets.
#sock_dir: /var/run/salt/minion

# In order to calculate the fqdns grain, all the IP addresses from the minion
# are processed with underlying calls to `socket.gethostbyaddr` which can take
# 5 seconds to be released (after reaching `socket.timeout`) when there is no
```

```
# fqdn for that IP. These calls to `socket.gethostbyaddr` are processed
# asynchronously, however, it still adds 5 seconds every time grains are
# generated if an IP does not resolve. In Windows grains are regenerated each
# time a new process is spawned. Therefore, the default for Windows is `False`.
# All other OSes default to `True`
# enable_fqdns_grains: True
```

```
# The minion can take a while to start up when lspci and/or dmidecode is used
# to populate the grains for the minion. Set this to False if you do not need
# GPU hardware grains for your minion.
# enable_gpu_grains: True
```

```
# Set the default outputter used by the salt-call command. The default is
# "nested".
#output: nested
```

```
# To set a list of additional directories to search for salt outputters, set the
# outputter_dirs option.
#outputter_dirs: []
```

```
# By default output is colored. To disable colored output, set the color value
# to False.
#color: True
```

```
# Do not strip off the colored output from nested results and state outputs
# (true by default).
# strip_colors: False
```

```
# Backup files that are replaced by file.managed and file.recurse under
```

'cachedir'/file_backup relative to their original location and appended
with a timestamp. The only valid setting is "minion". Disabled by default.

#

Alternatively this can be specified for each file in state files:

/etc/ssh/sshd_config:

file.managed:

- source: salt://ssh/sshd_config

- backup: minion

#

#backup_mode: minion

When waiting for a master to accept the minion's public key, salt will
continuously attempt to reconnect until successful. This is the time, in
seconds, between those reconnection attempts.

#acceptance_wait_time: 10

If this is nonzero, the time between reconnection attempts will increase by
acceptance_wait_time seconds per iteration, up to this maximum. If this is
set to zero, the time between reconnection attempts will stay constant.

#acceptance_wait_time_max: 0

If the master rejects the minion's public key, retry instead of exiting.

Rejected keys will be handled the same as waiting on acceptance.

#rejected_retry: False

When the master key changes, the minion will try to re-auth itself to receive
the new master key. In larger environments this can cause a SYN flood on the
master because all minions try to re-auth immediately. To prevent this and
have a minion wait for a random amount of time, use this optional parameter.

The wait-time will be a random number of seconds between 0 and the defined value.

#random_reauth_delay: 60

To avoid overloading a master when many minions startup at once, a randomized

delay may be set to tell the minions to wait before connecting to the master.

This value is the number of seconds to choose from for a random number. For

example, setting this value to 60 will choose a random number of seconds to delay

on startup between zero seconds and sixty seconds. Setting to '0' will disable

this feature.

#random_startup_delay: 0

When waiting for a master to accept the minion's public key, salt will

continuously attempt to reconnect until successful. This is the timeout value,

in seconds, for each individual attempt. After this timeout expires, the minion

will wait for acceptance_wait_time seconds before trying again. Unless your master

is under unusually heavy load, this should be left at the default.

#auth_timeout: 60

Number of consecutive SaltReqTimeoutError that are acceptable when trying to

authenticate.

#auth_tries: 7

The number of attempts to connect to a master before giving up.

Set this to -1 for unlimited attempts. This allows for a master to have

downtime and the minion to reconnect to it later when it comes back up.

In 'failover' mode, it is the number of attempts for each set of masters.

In this mode, it will cycle through the list of masters for each attempt.

#

This is different than auth_tries because auth_tries attempts to
retry auth attempts with a single master. auth_tries is under the
assumption that you can connect to the master but not gain
authorization from it. master_tries will still cycle through all
the masters in a given try, so it is appropriate if you expect
occasional downtime from the master(s).

#master_tries: 1

If authentication fails due to SaltReqTimeoutError during a ping_interval,
cause sub minion process to restart.

#auth_safemode: False

Ping Master to ensure connection is alive (minutes).

#ping_interval: 0

To auto recover minions if master changes IP address (DDNS)

auth_tries: 10

auth_safemode: False

ping_interval: 2

#

Minions won't know master is missing until a ping fails. After the ping fail,
the minion will attempt authentication and likely fails out and cause a restart.
When the minion restarts it will resolve the masters IP and attempt to reconnect.

If you don't have any problems with syn-floods, don't bother with the
three recon_* settings described below, just leave the defaults!

#

The ZeroMQ pull-socket that binds to the masters publishing interface tries
to reconnect immediately, if the socket is disconnected (for example if


```
# the master processes are restarted). In large setups this will have all
# minions reconnect immediately which might flood the master (the ZeroMQ-default
# is usually a 100ms delay). To prevent this, these three recon_* settings
# can be used.

# recon_default: the interval in milliseconds that the socket should wait before
#     trying to reconnect to the master (1000ms = 1 second)
#
# recon_max: the maximum time a socket should wait. each interval the time to wait
#     is calculated by doubling the previous time. if recon_max is reached,
#     it starts again at recon_default. Short example:
#
#     reconnect 1: the socket will wait 'recon_default' milliseconds
#     reconnect 2: 'recon_default' * 2
#     reconnect 3: ('recon_default' * 2) * 2
#     reconnect 4: value from previous interval * 2
#     reconnect 5: value from previous interval * 2
#     reconnect x: if value >= recon_max, it starts again with recon_default
#
# recon_randomize: generate a random wait time on minion start. The wait time will
#     be a random value between recon_default and recon_default +
#     recon_max. Having all minions reconnect with the same recon_default
#     and recon_max value kind of defeats the purpose of being able to
#     change these settings. If all minions have the same values and your
#     setup is quite large (several thousand minions), they will still
#     flood the master. The desired behavior is to have timeframe within
#     all minions try to reconnect.
#
# Example on how to use these settings. The goal: have all minions reconnect within a
# 60 second timeframe on a disconnect.
```

```
# recon_default: 1000
# recon_max: 59000
# recon_randomize: True
#
# Each minion will have a randomized reconnect value between 'recon_default'
# and 'recon_default + recon_max', which in this example means between 1000ms
# 60000ms (or between 1 and 60 seconds). The generated random-value will be
# doubled after each attempt to reconnect. Lets say the generated random
# value is 11 seconds (or 11000ms).
# reconnect 1: wait 11 seconds
# reconnect 2: wait 22 seconds
# reconnect 3: wait 33 seconds
# reconnect 4: wait 44 seconds
# reconnect 5: wait 55 seconds
# reconnect 6: wait time is bigger than 60 seconds (recon_default + recon_max)
# reconnect 7: wait 11 seconds
# reconnect 8: wait 22 seconds
# reconnect 9: wait 33 seconds
# reconnect x: etc.
#
# In a setup with ~6000 hosts these settings would average the reconnects
# to about 100 per second and all hosts would be reconnected within 60 seconds.
# recon_default: 100
# recon_max: 5000
# recon_randomize: False
#
#
# The loop_interval sets how long in seconds the minion will wait between
# evaluating the scheduler and running cleanup tasks. This defaults to 1
```

second on the minion scheduler.

#loop_interval: 1

Some installations choose to start all job returns in a cache or a returner

and forgo sending the results back to a master. In this workflow, jobs

are most often executed with --async from the Salt CLI and then results

are evaluated by examining job caches on the minions or any configured returners.

WARNING: Setting this to False will **disable** returns back to the master.

#pub_ret: True

The grains can be merged, instead of overridden, using this option.

This allows custom grains to defined different subvalues of a dictionary

grain. By default this feature is disabled, to enable set grains_deep_merge

to ``True``.

#grains_deep_merge: False

The grains_refresh_every setting allows for a minion to periodically check

its grains to see if they have changed and, if so, to inform the master

of the new grains. This operation is moderately expensive, therefore

care should be taken not to set this value too low.

#

Note: This value is expressed in __minutes__!

#

A value of 10 minutes is a reasonable default.

#

If the value is set to zero, this check is disabled.

#grains_refresh_every: 1

Cache grains on the minion. Default is False.

#grains_cache: False

Cache rendered pillar data on the minion. Default is False.

This may cause 'cachedir'/pillar to contain sensitive data that should be

protected accordingly.

#minion_pillar_cache: False

Grains cache expiration, in seconds. If the cache file is older than this

number of seconds then the grains cache will be dumped and fully re-populated

with fresh data. Defaults to 5 minutes. Will have no effect if 'grains_cache'

is not enabled.

grains_cache_expiration: 300

Determines whether or not the salt minion should run scheduled mine updates.

Defaults to "True". Set to "False" to disable the scheduled mine updates

(this essentially just does not add the mine update function to the minion's

scheduler).

#mine_enabled: True

Determines whether or not scheduled mine updates should be accompanied by a job

return for the job cache. Defaults to "False". Set to "True" to include job

returns in the job cache for mine updates.

#mine_return_job: False

Example functions that can be run via the mine facility

NO mine functions are established by default.

Note these can be defined in the minion's pillar as well.

#mine_functions:

```
# test.ping: []

# network.ip_addrs:

# interface: eth0

# cidr: '10.0.0.0/8'


# The number of minutes between mine updates.

#mine_interval: 60


# Windows platforms lack posix IPC and must rely on slower TCP based inter-
# process communications. ipc_mode is set to 'tcp' on such systems.

#ipc_mode: ipc


# Overwrite the default tcp ports used by the minion when ipc_mode is set to 'tcp'

#tcp_pub_port: 4510

#tcp_pull_port: 4511


# Passing very large events can cause the minion to consume large amounts of
# memory. This value tunes the maximum size of a message allowed onto the
# minion event bus. The value is expressed in bytes.

#max_event_size: 1048576


# When a minion starts up it sends a notification on the event bus with a tag
# that looks like this: `salt/minion/<minion_id>/start`. For historical reasons
# the minion also sends a similar event with an event tag like this:
# `minion_start`. This duplication can cause a lot of clutter on the event bus
# when there are many minions. Set `enable_legacy_startup_events: False` in the
# minion config to ensure only the `salt/minion/<minion_id>/start` events are
# sent. Beginning with the `Sodium` Salt release this option will default to
# `False`
```

```
#enable_legacy_startup_events: True
```

```
# To detect failed master(s) and fire events on connect/disconnect, set  
# master_alive_interval to the number of seconds to poll the masters for  
# connection events.
```

```
#
```

```
#master_alive_interval: 30
```

```
# The minion can include configuration from other files. To enable this,  
# pass a list of paths to this option. The paths can be either relative or  
# absolute; if relative, they are considered to be relative to the directory  
# the main minion configuration file lives in (this file). Paths can make use  
# of shell-style globbing. If no files are matched by a path passed to this  
# option then the minion will log a warning message.
```

```
#
```

```
# Include a config file from some other path:
```

```
# include: /etc/salt/extra_config
```

```
#
```

```
# Include config from several files and directories:
```

```
#include:
```

```
# - /etc/salt/extra_config
```

```
# - /etc/roles/webserver
```

```
# The syndic minion can verify that it is talking to the correct master via the  
# key fingerprint of the higher-level master with the "syndic_finger" config.
```

```
#syndic_finger: "
```

```
#
```

```
#
```

```
#
```

Minion module management

#####

Disable specific modules. This allows the admin to limit the level of
access the master has to the minion. The default here is the empty list,
below is an example of how this needs to be formatted in the config file

#disable_modules:

- cmdmod

- test

#disable_returners: []

This is the reverse of disable_modules. The default, like disable_modules, is the empty list,
but if this option is set to *anything* then *only* those modules will load.
Note that this is a very large hammer and it can be quite difficult to keep the minion working
the way you think it should since Salt uses many modules internally itself. At a bare minimum
you need the following enabled or else the minion won't start.

#whitelist_modules:

- cmdmod

- test

- config

Modules can be loaded from arbitrary paths. This enables the easy deployment
of third party modules. Modules for returners and minions can be loaded.
Specify a list of extra directories to search for minion modules and
returners. These paths must be fully qualified!

#module_dirs: []

#returner_dirs: []

#states_dirs: []

#render_dirs: []

#utils_dirs: []

```
#
# A module provider can be statically overwritten or extended for the minion
# via the providers option, in this case the default module will be
# overwritten by the specified module. In this example the pkg module will
# be provided by the yumpkg5 module instead of the system default.
#providers:
# pkg: yumpkg5
#
# Enable Cython modules searching and loading. (Default: False)
#cython_enable: False
#
# Specify a max size (in bytes) for modules on import. This feature is currently
# only supported on *nix operating systems and requires psutil.
# modules_max_memory: -1
```

```
##### State Management Settings #####
```

```
#####
```

```
# The default renderer to use in SLS files. This is configured as a
# pipe-delimited expression. For example, jinja|yaml will first run jinja
# templating on the SLS file, and then load the result as YAML. This syntax is
# documented in further depth at the following URL:
```

```
#
# https://docs.saltproject.io/en/latest/ref/renderers/#composing-renderers
```

```
#
# NOTE: The "shebang" prefix (e.g. "#!jinja|yaml") described in the
# documentation linked above is for use in an SLS file to override the default
# renderer, it should not be used when configuring the renderer here.
```

```
#
```



```
#renderer: jinja|yaml

#

# The failhard option tells the minions to stop immediately after the first
# failure detected in the state execution. Defaults to False.

#failhard: False

#

# Reload the modules prior to a highstate run.

#autoload_dynamic_modules: True

#

# clean_dynamic_modules keeps the dynamic modules on the minion in sync with
# the dynamic modules on the master, this means that if a dynamic module is
# not on the master it will be deleted from the minion. By default, this is
# enabled and can be disabled by changing this value to False.

#clean_dynamic_modules: True

#

# Renamed from ``environment`` to ``saltenv``. If ``environment`` is used,
# ``saltenv`` will take its value. If both are used, ``environment`` will be
# ignored and ``saltenv`` will be used.

# Normally the minion is not isolated to any single environment on the master
# when running states, but the environment can be isolated on the minion side
# by statically setting it. Remember that the recommended way to manage
# environments is to isolate via the top file.

#saltenv: None

#

# Isolates the pillar environment on the minion side. This functions the same
# as the environment setting, but for pillar instead of states.

#pillarenv: None

#

# Set this option to True to force the pillarenv to be the same as the
```

```
# effective saltenv when running states. Note that if pillarenv is specified,
# this option will be ignored.
#pillarenv_from_saltenv: False
#
# Set this option to 'True' to force a 'KeyError' to be raised whenever an
# attempt to retrieve a named value from pillar fails. When this option is set
# to 'False', the failed attempt returns an empty string. Default is 'False'.
#pillar_raise_on_missing: False
#
# If using the local file directory, then the state top file name needs to be
# defined, by default this is top.sls.
#state_top: top.sls
#
# Run states when the minion daemon starts. To enable, set startup_states to:
# 'highstate' -- Execute state.highstate
# 'sls' -- Read in the sls_list option and execute the named sls files
# 'top' -- Read top_file option and execute based on that file on the Master
#startup_states: ""
#
# List of states to run when the minion starts up if startup_states is 'sls':
#sls_list:
# - edit.vim
# - hyper
#
# List of grains to pass in start event when minion starts up:
#start_event_grains:
# - machine_id
# - uuid
#
```

```
# Top file to execute if startup_states is 'top':

#top_file: ""

# Automatically aggregate all states that have support for mod_aggregate by
# setting to True. Or pass a list of state module names to automatically
# aggregate just those types.
#
# state_aggregate:
# - pkg
#
#state_aggregate: False

# Disable requisites during state runs by specifying a single requisite
# or a list of requisites to disable.
#
# disabled_requisites: require_in
#
# disabled_requisites:
# - require
# - require_in

##### File Directory Settings #####
#####

# The Salt Minion can redirect all file server operations to a local directory,
# this allows for the same state tree that is on the master to be used if
# copied completely onto the minion. This is a literal copy of the settings on
# the master but used to reference a local directory on the minion.

# Set the file client. The client defaults to looking on the master server for
```

```
# files, but can be directed to look at the local file directory setting
# defined below by setting it to "local". Setting a local file_client runs the
# minion in masterless mode.
#file_client: remote
```

```
# The file directory works on environments passed to the minion, each environment
# can have multiple root directories, the subdirectories in the multiple file
# roots cannot match, otherwise the downloaded files will not be able to be
# reliably ensured. A base environment is required to house the top file.
```

```
# Example:
```

```
# file_roots:
```

```
# base:
```

```
# - /srv/salt/
```

```
# dev:
```

```
# - /srv/salt/dev/services
```

```
# - /srv/salt/dev/states
```

```
# prod:
```

```
# - /srv/salt/prod/services
```

```
# - /srv/salt/prod/states
```

```
#
```

```
#file_roots:
```

```
# base:
```

```
# - /srv/salt
```

```
# Uncomment the line below if you do not want the file_server to follow
```

```
# symlinks when walking the filesystem tree. This is set to True
```

```
# by default. Currently this only applies to the default roots
```

```
# fileserver_backend.
```

```
#fileserver_followsymlinks: False
```

#

Uncomment the line below if you do not want symlinks to be
treated as the files they are pointing to. By default this is set to
False. By uncommenting the line below, any detected symlink while listing
files on the Master will not be returned to the Minion.

#fileserver_ignoresymlinks: True

#

By default, the Salt fileserver recurses fully into all defined environments
to attempt to find files. To limit this behavior so that the fileserver only
traverses directories with SLS files and special Salt directories like _modules,
enable the option below. This might be useful for installations where a file root
has a very large number of files and performance is negatively impacted. Default
is False.

#fileserver_limit_traversal: False

The hash_type is the hash to use when discovering the hash of a file on
the local fileserver. The default is sha256, but md5, sha1, sha224, sha384
and sha512 are also supported.

#

WARNING: While md5 and sha1 are also supported, do not use them due to the
high chance of possible collisions and thus security breach.

#

Warning: Prior to changing this value, the minion should be stopped and all
Salt caches should be cleared.

#hash_type: sha256

The Salt pillar is searched for locally if file_client is set to local. If
this is the case, and pillar data is defined, then the pillar_roots need to
also be configured on the minion:

```
#pillar_roots:

# base:

# - /srv/pillar


# Set a hard-limit on the size of the files that can be pushed to the master.
# It will be interpreted as megabytes. Default: 100
#file_recv_max_size: 100

#

#

##### Security settings #####

#####

# Enable "open mode", this mode still maintains encryption, but turns off
# authentication, this is only intended for highly secure environments or for
# the situation where your keys end up in a bad state. If you run in open mode
# you do so at your own risk!
#open_mode: False


# The size of key that should be generated when creating new keys.
#keysize: 2048


# Enable permissive access to the salt keys. This allows you to run the
# master or minion as root, but have a non-root group be given access to
# your pki_dir. To make the access explicit, root must belong to the group
# you've given access to. This is potentially quite insecure.
#permissive_pki_access: False


# The state_verbose and state_output settings can be used to change the way
# state system data is printed to the display. By default all data is printed.
# The state_verbose setting can be set to True or False, when set to False
```

all data that has a result of True and no changes will be suppressed.

#state_verbose: True

The state_output setting controls which results will be output full multi line

full, terse - each state will be full/terse

mixed - only states with errors will be full

changes - states with changes and errors will be full

full_id, mixed_id, changes_id and terse_id are also allowed;

when set, the state ID will be used as name in the output

#state_output: full

The state_output_diff setting changes whether or not the output from

successful states is returned. Useful when even the terse output of these

states is cluttering the logs. Set it to True to ignore them.

#state_output_diff: False

The state_output_profile setting changes whether profile information

will be shown for each state run.

#state_output_profile: True

Fingerprint of the master public key to validate the identity of your Salt master

before the initial key exchange. The master fingerprint can be found by running

"salt-key -f master.pub" on the Salt master.

#master_finger: "

Use TLS/SSL encrypted connection between master and minion.

Can be set to a dictionary containing keyword arguments corresponding to Python's

'ssl.wrap_socket' method.

Default is None.

#ssl:

keyfile: <path_to_keyfile>

certfile: <path_to_certfile>

ssl_version: PROTOCOL_TLSv1_2

Grains to be sent to the master on authentication to check if the minion's key

will be accepted automatically. Needs to be configured on the master.

#autosign_grains:

- uuid

- server_id

Reactor Settings

#####

Define a salt reactor. See <https://docs.saltproject.io/en/latest/topics/reactor/>

#reactor: []

#Set the TTL for the cache of the reactor configuration.

#reactor_refresh_interval: 60

#Configure the number of workers for the runner/wheel in the reactor.

#reactor_worker_threads: 10

#Define the queue size for workers in the reactor.

#reactor_worker_hwm: 10000

Thread settings

#####


```
# Disable multiprocessing support, by default when a minion receives a
# publication a new process is spawned and the command is executed therein.
#
# WARNING: Disabling multiprocessing may result in substantial slowdowns
# when processing large pillars. See https://github.com/saltstack/salt/issues/38758
# for a full explanation.
#multiprocessing: True
```

```
# Limit the maximum amount of processes or threads created by salt-minion.
# This is useful to avoid resource exhaustion in case the minion receives more
# publications than it is able to handle, as it limits the number of spawned
# processes or threads. -1 is the default and disables the limit.
#process_count_max: -1
```

```
##### Logging settings #####
```

```
#####
```

```
# The location of the minion log file
# The minion log can be sent to a regular file, local path name, or network
# location. Remote logging works best when configured to use rsyslogd(8) (e.g.:
# ``file:///dev/log``), with rsyslogd(8) configured for network logging. The URI
# format is: <file|udp|tcp>://<host|socketpath>:<port-if-required>/<log-facility>
#log_file: /var/log/salt/minion
#log_file: file:///dev/log
#log_file: udp://loghost:10514
#
#log_file: /var/log/salt/minion
#key_logfile: /var/log/salt/key
```

```
# The level of messages to send to the console.

# One of 'garbage', 'trace', 'debug', 'info', 'warning', 'error', 'critical'.

#

# The following log levels are considered INSECURE and may log sensitive data:

# ['garbage', 'trace', 'debug']

#

# Default: 'warning'

#log_level: warning


# The level of messages to send to the log file.

# One of 'garbage', 'trace', 'debug', 'info', 'warning', 'error', 'critical'.

# If using 'log_granular_levels' this must be set to the highest desired level.

# Default: 'warning'

#log_level_logfile:


# The date and time format used in log messages. Allowed date/time formatting
# can be seen here: http://docs.python.org/library/time.html#time.strftime

#log_datefmt: '%H:%M:%S'

#log_datefmt_logfile: '%Y-%m-%d %H:%M:%S'


# The format of the console logging messages. Allowed formatting options can
# be seen here: http://docs.python.org/library/logging.html#logrecord-attributes

#

# Console log colors are specified by these additional formatters:

#

# %(colorlevel)s

# %(colorname)s

# %(colorprocess)s

# %(colormsg)s
```

```

#
# Since it is desirable to include the surrounding brackets, '[' and ']', in
# the coloring of the messages, these color formatters also include padding as
# well. Color LogRecord attributes are only available for console logging.
#
#log_fmt_console: '%(colorlevel)s %(colormsg)s'
#log_fmt_console: '[(levelname)-8s] %(message)s'
#
#log_fmt_logfile: '%(asctime)s,%(msecs)03d [(name)-17s] [(levelname)-8s] %(message)s'

# This can be used to control logging levels more specifically. This
# example sets the main salt library at the 'warning' level, but sets
# 'salt.modules' to log at the 'debug' level:
# log_granular_levels:
#   'salt': 'warning'
#   'salt.modules': 'debug'
#
#log_granular_levels: {}

# To diagnose issues with minions disconnecting or missing returns, ZeroMQ
# supports the use of monitor sockets to log connection events. This
# feature requires ZeroMQ 4.0 or higher.
#
# To enable ZeroMQ monitor sockets, set 'zmq_monitor' to 'True' and log at a
# debug level or higher.
#
# A sample log event is as follows:
#
# [DEBUG ] ZeroMQ event: {'endpoint': 'tcp://127.0.0.1:4505', 'event': 512,

```

```

# 'value': 27, 'description': 'EVENT_DISCONNECTED'}
#
# All events logged will include the string 'ZeroMQ event'. A connection event
# should be logged as the minion starts up and initially connects to the
# master. If not, check for debug log level and that the necessary version of
# ZeroMQ is installed.
#
#zmq_monitor: False

# Number of times to try to authenticate with the salt master when reconnecting
# to the master
#tcp_authentication_retries: 5

#####  Module configuration  #####
#####

# Salt allows for modules to be passed arbitrary configuration data, any data
# passed here in valid yaml format will be passed on to the salt minion modules
# for use. It is STRONGLY recommended that a naming convention be used in which
# the module name is followed by a . and then the value. Also, all top level
# data must be applied via the yaml dict construct, some examples:
#
# You can specify that all modules should run in test mode:
#test: True
#
# A simple value for the test module:
#test.foo: foo
#
# A list for the test module:
#test.bar: [baz,quo]

```

```

#
# A dict for the test module:
#test.baz: {spam: sausage, cheese: bread}
#
#
#####    Update settings    #####
#####

# Using the features in Esky, a salt minion can both run as a frozen app and
# be updated on the fly. These options control how the update process
# (saltutil.update()) behaves.
#
# The url for finding and downloading updates. Disabled by default.
#update_url: False
#
# The list of services to restart after a successful update. Empty by default.
#update_restart_services: []


#####    Keepalive settings    #####
#####

# ZeroMQ now includes support for configuring SO_KEEPALIVE if supported by
# the OS. If connections between the minion and the master pass through
# a state tracking device such as a firewall or VPN gateway, there is
# the risk that it could tear down the connection the master and minion
# without informing either party that their connection has been taken away.
# Enabling TCP Keepalives prevents this from happening.

# Overall state of TCP Keepalives, enable (1 or True), disable (0 or False)
# or leave to the OS defaults (-1), on Linux, typically disabled. Default True, enabled.

```

#tcp_keepalive: True

How long before the first keepalive should be sent in seconds. Default 300

to send the first keepalive after 5 minutes, OS default (-1) is typically 7200 seconds

on Linux see /proc/sys/net/ipv4/tcp_keepalive_time.

#tcp_keepalive_idle: 300

How many lost probes are needed to consider the connection lost. Default -1

to use OS defaults, typically 9 on Linux, see /proc/sys/net/ipv4/tcp_keepalive_probes.

#tcp_keepalive_cnt: -1

How often, in seconds, to send keepalives after the first one. Default -1 to

use OS defaults, typically 75 seconds on Linux, see

/proc/sys/net/ipv4/tcp_keepalive_intvl.

#tcp_keepalive_intvl: -1

Windows Software settings

#####

Location of the repository cache file on the master:

#win_repo_cache: 'salt://win/repo/winrepo.p'

Returner settings

#####

Default Minion returners. Can be a comma delimited string or a list:

#

#return: mysql

#

#return: mysql,slack,redis

#

#return:

- mysql

- hipchat

- slack

Miscellaneous settings

#####

Default match type for filtering events tags: startswith, endswith, find, regex, fnmatch

#event_match_type: startswith

EXAMPLE PROXY MINION CONFIGURATION FILE

Primary configuration settings

#####

This configuration file is used to manage the behavior of all Salt Proxy

Minions on this host.

With the exception of the location of the Salt Master Server, values that are

commented out but have an empty line after the comment are defaults that need

not be set in the config. If there is no blank line after the comment, the

value is presented as an example and is not the default.

Per default the proxy minion will automatically include all config files

from proxy.d/*.conf (proxy.d is a directory in the same directory

as the main minion config file).

#default_include: proxy.d/*.conf

Backwards compatibility option for proxymodules created before 2015.8.2

This setting will default to 'False' in the 2016.3.0 release

```
# Setting this to True adds proxymodules to the __opts__ dictionary.
# This breaks several Salt features (basically anything that serializes
# __opts__ over the wire) but retains backwards compatibility.
#add_proxymodule_to_opts: True

# Set the location of the salt master server. If the master server cannot be
# resolved, then the minion will fail to start.
#master: salt

# If a proxymodule has a function called 'grains', then call it during
# regular grains loading and merge the results with the proxy's grains
# dictionary. Otherwise it is assumed that the module calls the grains
# function in a custom way and returns the data elsewhere
#
# Default to False for 2016.3 and 2016.11. Switch to True for 2017.7.0.
# proxy_merge_grains_in_module: True

# If a proxymodule has a function called 'alive' returning a boolean
# flag reflecting the state of the connection with the remote device,
# when this option is set as True, a scheduled job on the proxy will
# try restarting the connection. The polling frequency depends on the
# next option, 'proxy_keep_alive_interval'. Added in 2017.7.0.
# proxy_keep_alive: True

# The polling interval (in minutes) to check if the underlying connection
# with the remote device is still alive. This option requires
# 'proxy_keep_alive' to be configured as True and the proxymodule to
# implement the 'alive' function. Added in 2017.7.0.
# proxy_keep_alive_interval: 1
```


By default, any proxy opens the connection with the remote device when
initialized. Some proxymodules allow through this option to open/close
the session per command. This requires the proxymodule to have this
capability. Please consult the documentation to see if the proxy type
used can be that flexible. Added in 2017.7.0.
proxy_always_alive: True

If multiple masters are specified in the 'master' setting, the default behavior
is to always try to connect to them in the order they are listed. If random_master is
set to True, the order will be randomized instead. This can be helpful in distributing
the load of many minions executing salt-call requests, for example, from a cron job.
If only one master is listed, this setting is ignored and a warning will be logged.
#random_master: False

Minions can connect to multiple masters simultaneously (all masters
are "hot"), or can be configured to failover if a master becomes
unavailable. Multiple hot masters are configured by setting this
value to "str". Failover masters can be requested by setting
to "failover". MAKE SURE TO SET master_alive_interval if you are
using failover.
master_type: str

Poll interval in seconds for checking if the master is still there. Only
respected if master_type above is "failover".
master_alive_interval: 30

Set whether the minion should connect to the master via IPv6:
#ipv6: False

```
# Set the number of seconds to wait before attempting to resolve
# the master hostname if name resolution fails. Defaults to 30 seconds.
# Set to zero if the minion should shutdown and not retry.
# retry_dns: 30

# Set the port used by the master reply and authentication server.
#master_port: 4506

# The user to run salt.
#user: root

# Setting sudo_user will cause salt to run all execution modules under an sudo
# to the user given in sudo_user. The user under which the salt minion process
# itself runs will still be that provided in the user config above, but all
# execution modules run by the minion will be rerouted through sudo.
#sudo_user: saltdev

# Specify the location of the daemon process ID file.
#pidfile: /var/run/salt-minion.pid

# The root directory prepended to these options: pki_dir, cachedir, log_file,
# sock_dir, pidfile.
#root_dir: /

# The directory to store the pki information in
#pki_dir: /etc/salt/pki/minion

# Where cache data goes.
```

This data may contain sensitive data and should be protected accordingly.

#cachedir: /var/cache/salt/minion

Append minion_id to these directories. Helps with

multiple proxies and minions running on the same machine.

Allowed elements in the list: pki_dir, cachedir, extension_modules

Normally not needed unless running several proxies and/or minions on the same machine

Defaults to ['cachedir'] for proxies, [] (empty list) for regular minions

append_minionid_config_dirs:

- cachedir

Verify and set permissions on configuration directories at startup.

#verify_env: True

The minion can locally cache the return data from jobs sent to it, this

can be a good way to keep track of jobs the minion has executed

(on the minion side). By default this feature is disabled, to enable, set

cache_jobs to True.

#cache_jobs: False

Set the directory used to hold unix sockets.

#sock_dir: /var/run/salt/minion

Set the default outputter used by the salt-call command. The default is

"nested".

#output: nested

#

By default output is colored. To disable colored output, set the color value
to False.

#color: True

Do not strip off the colored output from nested results and state outputs
(true by default).

strip_colors: False

Backup files that are replaced by file.managed and file.recurse under
'cachedir'/file_backup relative to their original location and appended
with a timestamp. The only valid setting is "minion". Disabled by default.

#

Alternatively this can be specified for each file in state files:

/etc/ssh/sshd_config:

file.managed:

- source: salt://ssh/sshd_config

- backup: minion

#

#backup_mode: minion

When waiting for a master to accept the minion's public key, salt will
continuously attempt to reconnect until successful. This is the time, in
seconds, between those reconnection attempts.

#acceptance_wait_time: 10

If this is nonzero, the time between reconnection attempts will increase by
acceptance_wait_time seconds per iteration, up to this maximum. If this is
set to zero, the time between reconnection attempts will stay constant.

#acceptance_wait_time_max: 0

If the master rejects the minion's public key, retry instead of exiting.

Rejected keys will be handled the same as waiting on acceptance.

#rejected_retry: False

When the master key changes, the minion will try to re-auth itself to receive

the new master key. In larger environments this can cause a SYN flood on the

master because all minions try to re-auth immediately. To prevent this and

have a minion wait for a random amount of time, use this optional parameter.

The wait-time will be a random number of seconds between 0 and the defined value.

#random_reauth_delay: 60

When waiting for a master to accept the minion's public key, salt will

continuously attempt to reconnect until successful. This is the timeout value,

in seconds, for each individual attempt. After this timeout expires, the minion

will wait for acceptance_wait_time seconds before trying again. Unless your master

is under unusually heavy load, this should be left at the default.

#auth_timeout: 60

Number of consecutive SaltReqTimeoutError that are acceptable when trying to

authenticate.

#auth_tries: 7

If authentication fails due to SaltReqTimeoutError during a ping_interval,

cause sub minion process to restart.

#auth_safemode: False

Ping Master to ensure connection is alive (minutes).

#ping_interval: 0

```
# To auto recover minions if master changes IP address (DDNS)

# auth_tries: 10

# auth_safemode: False

# ping_interval: 90

#

# Minions won't know master is missing until a ping fails. After the ping fail,
# the minion will attempt authentication and likely fails out and cause a restart.
# When the minion restarts it will resolve the masters IP and attempt to reconnect.


# If you don't have any problems with syn-floods, don't bother with the
# three recon_* settings described below, just leave the defaults!

#

# The ZeroMQ pull-socket that binds to the masters publishing interface tries
# to reconnect immediately, if the socket is disconnected (for example if
# the master processes are restarted). In large setups this will have all
# minions reconnect immediately which might flood the master (the ZeroMQ-default
# is usually a 100ms delay). To prevent this, these three recon_* settings
# can be used.

# recon_default: the interval in milliseconds that the socket should wait before
#         trying to reconnect to the master (1000ms = 1 second)

#

# recon_max: the maximum time a socket should wait. each interval the time to wait
#         is calculated by doubling the previous time. if recon_max is reached,
#         it starts again at recon_default. Short example:

#

#         reconnect 1: the socket will wait 'recon_default' milliseconds
#         reconnect 2: 'recon_default' * 2
#         reconnect 3: ('recon_default' * 2) * 2
```

```
#      reconnect 4: value from previous interval * 2
#      reconnect 5: value from previous interval * 2
#      reconnect x: if value >= recon_max, it starts again with recon_default
#
# recon_randomize: generate a random wait time on minion start. The wait time will
#                  be a random value between recon_default and recon_default +
#                  recon_max. Having all minions reconnect with the same recon_default
#                  and recon_max value kind of defeats the purpose of being able to
#                  change these settings. If all minions have the same values and your
#                  setup is quite large (several thousand minions), they will still
#                  flood the master. The desired behavior is to have timeframe within
#                  all minions try to reconnect.
#
# Example on how to use these settings. The goal: have all minions reconnect within a
# 60 second timeframe on a disconnect.
# recon_default: 1000
# recon_max: 59000
# recon_randomize: True
#
# Each minion will have a randomized reconnect value between 'recon_default'
# and 'recon_default + recon_max', which in this example means between 1000ms
# 60000ms (or between 1 and 60 seconds). The generated random-value will be
# doubled after each attempt to reconnect. Lets say the generated random
# value is 11 seconds (or 11000ms).
# reconnect 1: wait 11 seconds
# reconnect 2: wait 22 seconds
# reconnect 3: wait 33 seconds
# reconnect 4: wait 44 seconds
# reconnect 5: wait 55 seconds
```

```
# reconnect 6: wait time is bigger than 60 seconds (recon_default + recon_max)

# reconnect 7: wait 11 seconds

# reconnect 8: wait 22 seconds

# reconnect 9: wait 33 seconds

# reconnect x: etc.

#

# In a setup with ~6000 thousand hosts these settings would average the reconnects

# to about 100 per second and all hosts would be reconnected within 60 seconds.

# recon_default: 100

# recon_max: 5000

# recon_randomize: False

#

#

# The loop_interval sets how long in seconds the minion will wait between

# evaluating the scheduler and running cleanup tasks. This defaults to a

# sane 60 seconds, but if the minion scheduler needs to be evaluated more

# often lower this value

#loop_interval: 60


# The grains_refresh_every setting allows for a minion to periodically check

# its grains to see if they have changed and, if so, to inform the master

# of the new grains. This operation is moderately expensive, therefore

# care should be taken not to set this value too low.

#

# Note: This value is expressed in __minutes__!

#

# A value of 10 minutes is a reasonable default.

#

# If the value is set to zero, this check is disabled.
```


#grains_refresh_every: 1

Cache grains on the minion. Default is False.

#grains_cache: False

Grains cache expiration, in seconds. If the cache file is older than this

number of seconds then the grains cache will be dumped and fully re-populated

with fresh data. Defaults to 5 minutes. Will have no effect if 'grains_cache'

is not enabled.

grains_cache_expiration: 300

Windows platforms lack posix IPC and must rely on slower TCP based inter-

process communications. Set ipc_mode to 'tcp' on such systems

#ipc_mode: ipc

Overwrite the default tcp ports used by the minion when in tcp mode

#tcp_pub_port: 4510

#tcp_pull_port: 4511

Passing very large events can cause the minion to consume large amounts of

memory. This value tunes the maximum size of a message allowed onto the

minion event bus. The value is expressed in bytes.

#max_event_size: 1048576

To detect failed master(s) and fire events on connect/disconnect, set

master_alive_interval to the number of seconds to poll the masters for

connection events.

#

#master_alive_interval: 30

```
# The minion can include configuration from other files. To enable this,
# pass a list of paths to this option. The paths can be either relative or
# absolute; if relative, they are considered to be relative to the directory
# the main minion configuration file lives in (this file). Paths can make use
# of shell-style globbing. If no files are matched by a path passed to this
# option then the minion will log a warning message.
#
# Include a config file from some other path:
# include: /etc/salt/extra_config
#
# Include config from several files and directories:
#include:
# - /etc/salt/extra_config
# - /etc/roles/webserver
#
#
#
##### Minion module management #####
#####
# Disable specific modules. This allows the admin to limit the level of
# access the master has to the minion.
#disable_modules: [cmd,test]
#disable_returners: []
#
# Modules can be loaded from arbitrary paths. This enables the easy deployment
# of third party modules. Modules for returners and minions can be loaded.
# Specify a list of extra directories to search for minion modules and
# returners. These paths must be fully qualified!
```

```

#module_dirs: []
#returner_dirs: []
#states_dirs: []
#render_dirs: []
#utils_dirs: []
#
# A module provider can be statically overwritten or extended for the minion
# via the providers option, in this case the default module will be
# overwritten by the specified module. In this example the pkg module will
# be provided by the yumpkg5 module instead of the system default.
#providers:
# pkg: yumpkg5
#
# Enable Cython modules searching and loading. (Default: False)
#cython_enable: False
#
# Specify a max size (in bytes) for modules on import. This feature is currently
# only supported on *nix operating systems and requires psutil.
# modules_max_memory: -1

##### State Management Settings #####
#####

# The default renderer to use in SLS files. This is configured as a
# pipe-delimited expression. For example, jinja|yaml will first run jinja
# templating on the SLS file, and then load the result as YAML. This syntax is
# documented in further depth at the following URL:
#
# https://docs.saltproject.io/en/latest/ref/renderers/#composing-renderers

```

```
#
# NOTE: The "shebang" prefix (e.g. "#!jinja|yaml") described in the
# documentation linked above is for use in an SLS file to override the default
# renderer, it should not be used when configuring the renderer here.
#
#renderer: jinja|yaml
#
# The failhard option tells the minions to stop immediately after the first
# failure detected in the state execution. Defaults to False.
#failhard: False
#
# Reload the modules prior to a highstate run.
#autoload_dynamic_modules: True
#
# clean_dynamic_modules keeps the dynamic modules on the minion in sync with
# the dynamic modules on the master, this means that if a dynamic module is
# not on the master it will be deleted from the minion. By default, this is
# enabled and can be disabled by changing this value to False.
#clean_dynamic_modules: True
#
# Normally, the minion is not isolated to any single environment on the master
# when running states, but the environment can be isolated on the minion side
# by statically setting it. Remember that the recommended way to manage
# environments is to isolate via the top file.
#environment: None
#
# If using the local file directory, then the state top file name needs to be
# defined, by default this is top.sls.
#state_top: top.sls
```

```

#

# Run states when the minion daemon starts. To enable, set startup_states to:

# 'highstate' -- Execute state.highstate

# 'sls' -- Read in the sls_list option and execute the named sls files

# 'top' -- Read top_file option and execute based on that file on the Master

#startup_states: "

#

# List of states to run when the minion starts up if startup_states is 'sls':

#sls_list:

# - edit.vim

# - hyper

#

# Top file to execute if startup_states is 'top':

#top_file: "


# Automatically aggregate all states that have support for mod_aggregate by
# setting to True. Or pass a list of state module names to automatically
# aggregate just those types.

#

# state_aggregate:

# - pkg

#

#state_aggregate: False


##### File Directory Settings #####

#####

# The Salt Minion can redirect all file server operations to a local directory,
# this allows for the same state tree that is on the master to be used if
# copied completely onto the minion. This is a literal copy of the settings on

```

the master but used to reference a local directory on the minion.

Set the file client. The client defaults to looking on the master server for
files, but can be directed to look at the local file directory setting
defined below by setting it to "local". Setting a local file_client runs the
minion in masterless mode.

#file_client: remote

The file directory works on environments passed to the minion, each environment
can have multiple root directories, the subdirectories in the multiple file
roots cannot match, otherwise the downloaded files will not be able to be
reliably ensured. A base environment is required to house the top file.

Example:

file_roots:

base:

- /srv/salt/

dev:

- /srv/salt/dev/services

- /srv/salt/dev/states

prod:

- /srv/salt/prod/services

- /srv/salt/prod/states

#

#file_roots:

base:

- /srv/salt

By default, the Salt fileserver recurses fully into all defined environments
to attempt to find files. To limit this behavior so that the fileserver only

traverses directories with SLS files and special Salt directories like _modules,
enable the option below. This might be useful for installations where a file root
has a very large number of files and performance is negatively impacted. Default
is False.

#fileserver_limit_traversal: False

The hash_type is the hash to use when discovering the hash of a file in
the local fileserver. The default is sha256 but sha224, sha384 and sha512
are also supported.

#

WARNING: While md5 and sha1 are also supported, do not use it due to the high chance
of possible collisions and thus security breach.

#

WARNING: While md5 is also supported, do not use it due to the high chance
of possible collisions and thus security breach.

#

Warning: Prior to changing this value, the minion should be stopped and all
Salt caches should be cleared.

#hash_type: sha256

The Salt pillar is searched for locally if file_client is set to local. If
this is the case, and pillar data is defined, then the pillar_roots need to
also be configured on the minion:

#pillar_roots:

base:

- /srv/pillar

#

#

Security settings

```
#####
```

```
# Enable "open mode", this mode still maintains encryption, but turns off
# authentication, this is only intended for highly secure environments or for
# the situation where your keys end up in a bad state. If you run in open mode
# you do so at your own risk!
#open_mode: False
```

```
# Enable permissive access to the salt keys. This allows you to run the
# master or minion as root, but have a non-root group be given access to
# your pki_dir. To make the access explicit, root must belong to the group
# you've given access to. This is potentially quite insecure.
#permissive_pki_access: False
```

```
# The state_verbose and state_output settings can be used to change the way
# state system data is printed to the display. By default all data is printed.
# The state_verbose setting can be set to True or False, when set to False
# all data that has a result of True and no changes will be suppressed.
#state_verbose: True
```

```
# The state_output setting controls which results will be output full multi line
# full, terse - each state will be full/terse
# mixed - only states with errors will be full
# changes - states with changes and errors will be full
# full_id, mixed_id, changes_id and terse_id are also allowed;
# when set, the state ID will be used as name in the output
#state_output: full
```

```
# The state_output_diff setting changes whether or not the output from
# successful states is returned. Useful when even the terse output of these
```


states is cluttering the logs. Set it to True to ignore them.

#state_output_diff: False

The state_output_profile setting changes whether profile information

will be shown for each state run.

#state_output_profile: True

Fingerprint of the master public key to validate the identity of your Salt master

before the initial key exchange. The master fingerprint can be found by running

"salt-key -F master" on the Salt master.

#master_finger: "

Thread settings

#####

Disable multiprocessing support, by default when a minion receives a

publication a new process is spawned and the command is executed therein.

#multiprocessing: True

Logging settings

#####

The location of the minion log file

The minion log can be sent to a regular file, local path name, or network

location. Remote logging works best when configured to use rsyslogd(8) (e.g.:

``file:///dev/log``), with rsyslogd(8) configured for network logging. The URI

format is: <file|udp|tcp>://<host|socketpath>:<port-if-required>/<log-facility>

#log_file: /var/log/salt/minion

#log_file: file:///dev/log

#log_file: udp://loghost:10514

#

#log_file: /var/log/salt/minion

#key_logfile: /var/log/salt/key

The level of messages to send to the console.

One of 'garbage', 'trace', 'debug', 'info', 'warning', 'error', 'critical'.

#

The following log levels are considered INSECURE and may log sensitive data:

['garbage', 'trace', 'debug']

#

Default: 'warning'

#log_level: warning

The level of messages to send to the log file.

One of 'garbage', 'trace', 'debug', 'info', 'warning', 'error', 'critical'.

If using 'log_granular_levels' this must be set to the highest desired level.

Default: 'warning'

#log_level_logfile:

The date and time format used in log messages. Allowed date/time formatting

can be seen here: <http://docs.python.org/library/time.html#time.strftime>

#log_datefmt: '%H:%M:%S'

#log_datefmt_logfile: '%Y-%m-%d %H:%M:%S'

The format of the console logging messages. Allowed formatting options can

be seen here: <http://docs.python.org/library/logging.html#logrecord-attributes>

#

Console log colors are specified by these additional formatters:

```

#
# %(colorlevel)s
# %(colorname)s
# %(colorprocess)s
# %(colormsg)s
#
# Since it is desirable to include the surrounding brackets, '[' and ']', in
# the coloring of the messages, these color formatters also include padding as
# well. Color LogRecord attributes are only available for console logging.
#
#log_fmt_console: '%(colorlevel)s %(colormsg)s'
#log_fmt_console: '[(levelname)-8s] %(message)s'
#
#log_fmt_logfile: '%(asctime)s,%(msecs)03d [(name)-17s] [(levelname)-8s] %(message)s'

# This can be used to control logging levels more specifically. This
# example sets the main salt library at the 'warning' level, but sets
# 'salt.modules' to log at the 'debug' level:
# log_granular_levels:
#   'salt': 'warning'
#   'salt.modules': 'debug'
#
#log_granular_levels: {}

# To diagnose issues with minions disconnecting or missing returns, ZeroMQ
# supports the use of monitor sockets # to log connection events. This
# feature requires ZeroMQ 4.0 or higher.
#
# To enable ZeroMQ monitor sockets, set 'zmq_monitor' to 'True' and log at a

```

```

# debug level or higher.

#
# A sample log event is as follows:
#
# [DEBUG ] ZeroMQ event: {'endpoint': 'tcp://127.0.0.1:4505', 'event': 512,
# 'value': 27, 'description': 'EVENT_DISCONNECTED'}
#
# All events logged will include the string 'ZeroMQ event'. A connection event
# should be logged on the as the minion starts up and initially connects to the
# master. If not, check for debug log level and that the necessary version of
# ZeroMQ is installed.
#
#zmq_monitor: False

#####  Module configuration  #####
#####

# Salt allows for modules to be passed arbitrary configuration data, any data
# passed here in valid yaml format will be passed on to the salt minion modules
# for use. It is STRONGLY recommended that a naming convention be used in which
# the module name is followed by a . and then the value. Also, all top level
# data must be applied via the yaml dict construct, some examples:
#
# You can specify that all modules should run in test mode:
#test: True
#
# A simple value for the test module:
#test.foo: foo
#
# A list for the test module:

```

```
#test.bar: [baz,quo]

#
# A dict for the test module:
#test.baz: {spam: sausage, cheese: bread}

#
#
##### Update settings #####
#####

# Using the features in Esky, a salt minion can both run as a frozen app and
# be updated on the fly. These options control how the update process
# (saltutil.update()) behaves.
#
# The url for finding and downloading updates. Disabled by default.
#update_url: False
#
# The list of services to restart after a successful update. Empty by default.
#update_restart_services: []


##### Keepalive settings #####
#####

# ZeroMQ now includes support for configuring SO_KEEPALIVE if supported by
# the OS. If connections between the minion and the master pass through
# a state tracking device such as a firewall or VPN gateway, there is
# the risk that it could tear down the connection the master and minion
# without informing either party that their connection has been taken away.
# Enabling TCP Keepalives prevents this from happening.

# Overall state of TCP Keepalives, enable (1 or True), disable (0 or False)
```

or leave to the OS defaults (-1), on Linux, typically disabled. Default True, enabled.

#tcp_keepalive: True

How long before the first keepalive should be sent in seconds. Default 300

to send the first keepalive after 5 minutes, OS default (-1) is typically 7200 seconds

on Linux see /proc/sys/net/ipv4/tcp_keepalive_time.

#tcp_keepalive_idle: 300

How many lost probes are needed to consider the connection lost. Default -1

to use OS defaults, typically 9 on Linux, see /proc/sys/net/ipv4/tcp_keepalive_probes.

#tcp_keepalive_cnt: -1

How often, in seconds, to send keepalives after the first one. Default -1 to

use OS defaults, typically 75 seconds on Linux, see

/proc/sys/net/ipv4/tcp_keepalive_intvl.

#tcp_keepalive_intvl: -1

Windows Software settings

#####

Location of the repository cache file on the master:

#win_repo_cache: 'salt://win/repo/winrepo.p'

Returner settings

#####

Which returner(s) will be used for minion's result:

#return: mysql